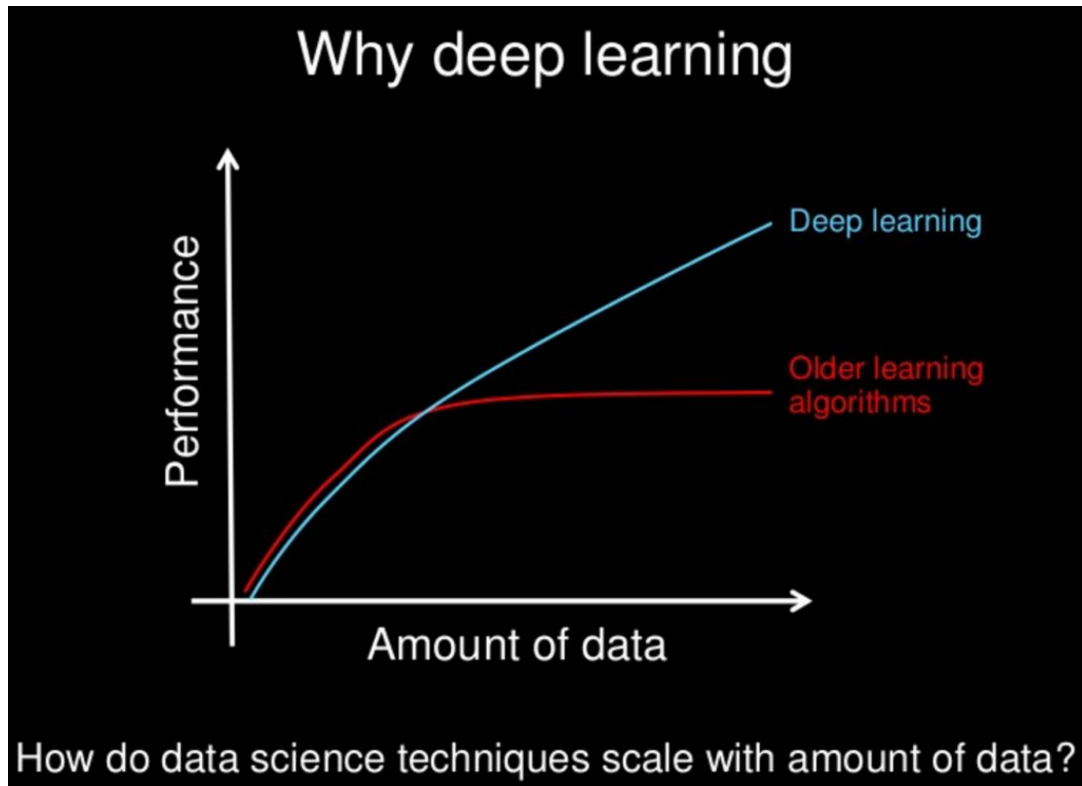


Uber's Distributed Deep Learning Journey




Alex Sergeev, Machine Learning Platform, Uber Engineering
@alsrgv

Deep Learning




Deep Learning @ Uber

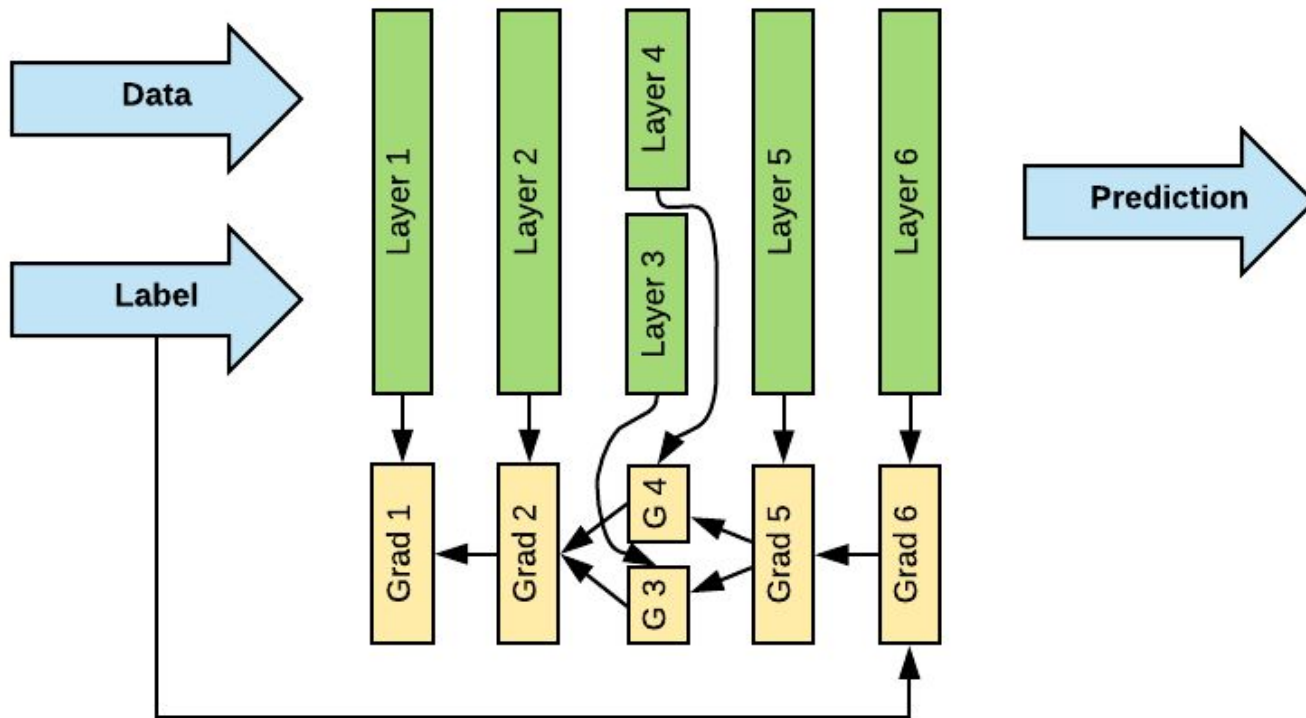
- Self-Driving Vehicles
- Trip Forecasting
- Fraud Detection
- ... and much more!



How does Deep Learning work?



How does Deep Learning training work?




TensorFlow

- Most popular open source framework for deep learning
- Combines high performance with ability to tinker with low level model details
- Has end-to-end support from research to production


Going Distributed

- Train very large models
- Speed up model training



Model Parallelism

VS

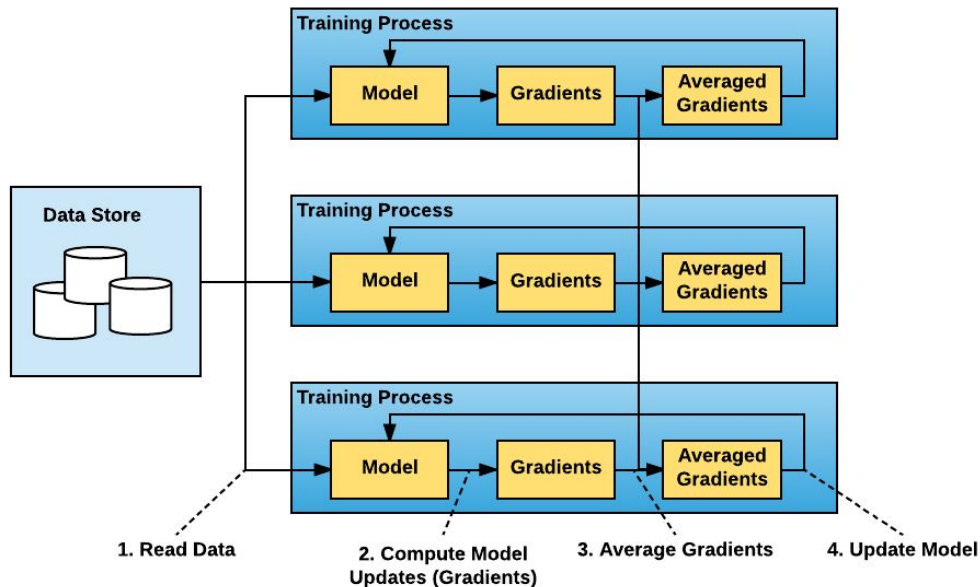


Data Parallelism

Going Distributed Cont.

- Modern GPUs have a lot of RAM
- Vast majority of use cases are data-parallel
- Facebook demonstrated training ResNet-50 on ImageNet in 1 hour

(arxiv.org/abs/1706.02677)



Parameter Server Technique

`tf.Server()`

`tf.train.replicas_device_setter()`

Parameter Server


Worker

GPU Towers

`tf.ClusterSpec()`


`tf.train.SyncReplicasOptimizer()`

Averages All the Gradients



Each Averages Portion of the Gradients

or



UBER

UBER | DATA

Parameter Server Technique - Example Script

```
import argparse
import sys
import tensorflow as tf

FLAGS = None

def main():
    ps_hosts = FLAGS.ps_hosts.split(",")
    worker_hosts = FLAGS.worker_hosts.split(",")

    # Create a cluster from the parameter server and worker hosts.
    cluster = tf.train.ClusterSpec({'ps': ps_hosts, 'worker': worker_hosts})

    # Create and start a server for the local task.
    server = tf.train.Server(cluster,
                             job_name=FLAGS.job_name,
                             task_index=FLAGS.task_index)

    if FLAGS.job_name == "ps":
        server.join()
    elif FLAGS.job_name == "worker":

        # Assigns ops to the local worker by default.
        with tf.device(tf.train.replica_device_setter(
            worker_device="/job:worker/task%d" % FLAGS.task_index,
            cluster=cluster)):

            # Build model...
            loss = ...
            global_step = tf.contrib.framework.get_or_create_global_step()
            train_op = tf.train.AdamOptimizer(0.01).minimize(
                loss, global_step=global_step)


            # The StopAtStepHook handles stopping after running given steps.
            hooks=[tf.train.StopAtStepHook(last_step=1000000)]

            # The MonitoredTrainingSession takes care of session initialization,
            # restoring from a checkpoint, saving to a checkpoint, and closing when done
            # or an error occurs.
            with tf.train.MonitoredTrainingSession(master=server.target,
                                                    is_chief=(FLAGS.task_index == 0),
                                                    checkpoint_dir="/tmp/train_logs",
                                                    hooks=hooks) as mon_sess:

                while not mon_sess.should_stop():
                    # Run a training step asynchronously.
                    # See "tf.train.SyncReplicasOptimizer" for additional details on how to
                    # perform "synchronous" training.
                    # mon_sess.run handles AbortedError in case of preempted PS.
                    mon_sess.run(train_op)

if __name__ == '__main__':
    parser = argparse.ArgumentParser()
    parser.register('type', 'bool', lambda v: v.lower() == "true")
    # Flags for defining the tf.train.ClusterSpec
    parser.add_argument(
        "--ps_hosts",
        type=str,
        default="",
        help="Comma-separated list of hostname:port pairs"
    )
    parser.add_argument(
        "--worker_hosts",
        type=str,
        default="",
        help="Comma-separated list of hostname:port pairs"
    )
    parser.add_argument(
        "--job_name",
        type=str,
        default="One of 'ps', 'worker'"
    )
    # Flags for defining the tf.train.Server
    parser.add_argument(
        "--task_index",
        type=int,
        default=0,
        help="Index of task within the job"
    )
    FLAGS, unparsed = parser.parse_known_args()
```

Parameter Server Technique - Performance



UBER

Considering ImageNet dataset of 1.3M images, this allows to train ResNet-101 for one epoch in 3.5 minutes. Scaling efficiency on 128 GPUs is only 42%, however.

How Can We Improve?


- Re-think necessary complexity for data-parallel case
- Improve communication algorithm
- Use RDMA-capable networking (InfiniBand, RoCE)

Meet Horovod



- Distributed training framework for TensorFlow
- Inspired by HPC techniques and work of Baidu, Facebook, et al.
- Uses bandwidth-optimal communication protocols
 - Makes use of RDMA (InfiniBand, RoCE) if available
- Seamlessly installs on top of TensorFlow via
`pip install horovod`
- Named after traditional Russian folk dance where participants dance in a circle with linked hands

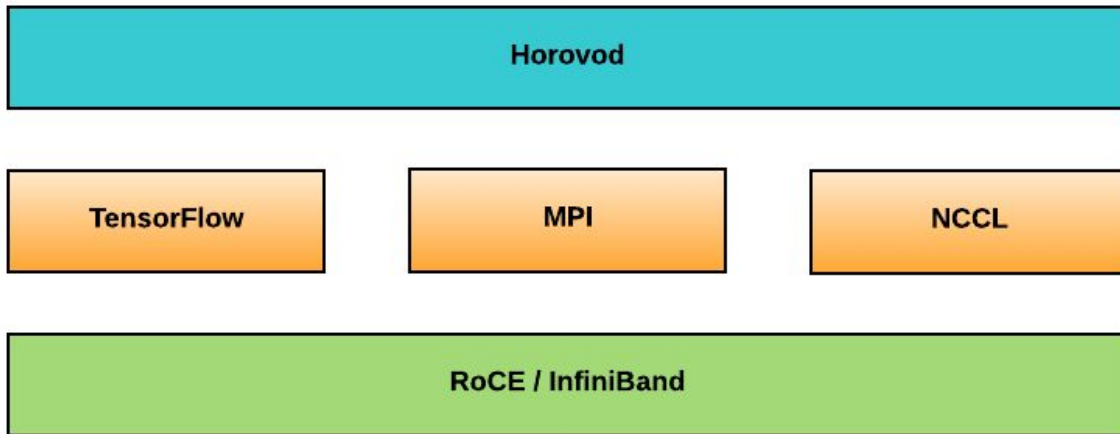
Horovod Technique



Patarasuk, P., & Yuan, X. (2009). Bandwidth optimal all-reduce algorithms for clusters of workstations. *Journal of Parallel and Distributed Computing*, 69(2), 117-124. doi:10.1016/j.jpdc.2008.09.002

Horovod Stack

- Plugs into TensorFlow via custom op mechanism
- Uses MPI for worker discovery and reduction coordination
- Uses NVIDIA NCCL for actual reduction on the server and across servers



Horovod Example

```
import tensorflow as tf
import horovod.tensorflow as hvd

# Initialize Horovod
hvd.init()

# Pin GPU to be used
config = tf.ConfigProto()
config.gpu_options.visible_device_list = str(hvd.local_rank())

# Build model...
loss = ...
opt = tf.train.AdagradOptimizer(0.01)

# Add Horovod Distributed Optimizer
opt = hvd.DistributedOptimizer(opt)

# Add hook to broadcast variables from rank 0 to all other processes during initialization.
hooks = [hvd.BroadcastGlobalVariablesHook(0)]

# Make training operation
train_op = opt.minimize(loss)

# The MonitoredTrainingSession takes care of session initialization,
# restoring from a checkpoint, saving to a checkpoint, and closing when done
# or an error occurs.
with tf.train.MonitoredTrainingSession(checkpoint_dir="/tmp/train_logs",
                                       config=config, hooks=hooks) as mon_sess:
    while not mon_sess.should_stop():
        # Perform synchronous training.
        mon_sess.run(train_op)
```


Horovod Example - Keras

```
import keras
from keras import backend as K
import tensorflow as tf
import horovod.keras as hvd

# Initialize Horovod.
hvd.init()

# Pin GPU to be used to process local rank (one GPU per process)
config = tf.ConfigProto()
config.gpu_options.allow_growth = True
config.gpu_options.visible_device_list = str(hvd.local_rank())
K.set_session(tf.Session(config=config))

# Build model...
model = ...
opt = keras.optimizers.Adadelta(1.0)

# Add Horovod Distributed Optimizer.
opt = hvd.DistributedOptimizer(opt)

model.compile(loss=keras.losses.categorical_crossentropy, optimizer=opt, metrics=['accuracy'])

# Broadcast initial variable states from rank 0 to all other processes.
callbacks = [hvd.callbacks.BroadcastGlobalVariablesCallback(0)]

model.fit(x_train, y_train,
        callbacks=callbacks,
        epochs=10,
        validation_data=(x_test, y_test))
```

Horovod Example - Estimator API

```
import tensorflow as tf
import horovod.tensorflow as hvd

# Initialize Horovod
hvd.init()

# Pin GPU to be used
config = tf.ConfigProto()
config.gpu_options.visible_device_list = str(hvd.local_rank())

# Build model...
def model_fn(features, labels, mode):
    loss = ...
    opt = tf.train.AdagradOptimizer(0.01)

# Add Horovod Distributed Optimizer
opt = hvd.DistributedOptimizer(opt)

train_op = optimizer.minimize(loss=loss, global_step=tf.train.get_global_step())
return tf.estimator.EstimatorSpec(mode=mode, loss=loss, train_op=train_op)

# Add hook to broadcast variables from rank 0 to all other processes during initialization.
hooks = [hvd.BroadcastGlobalVariablesHook(0)]

# Create the Estimator
mnist_classifier = tf.estimator.Estimator(
    model_fn=cnn_model_fn, model_dir="/tmp/mnist_convnet_model",
    config=tf.estimator.RunConfig(session_config=config))

mnist_classifier.train(input_fn=train_input_fn, steps=100, hooks=hooks)
```

Running Horovod

- MPI takes care of launching processes on all machines
- Run on a 4 GPU machine (Open MPI 3.0.0):

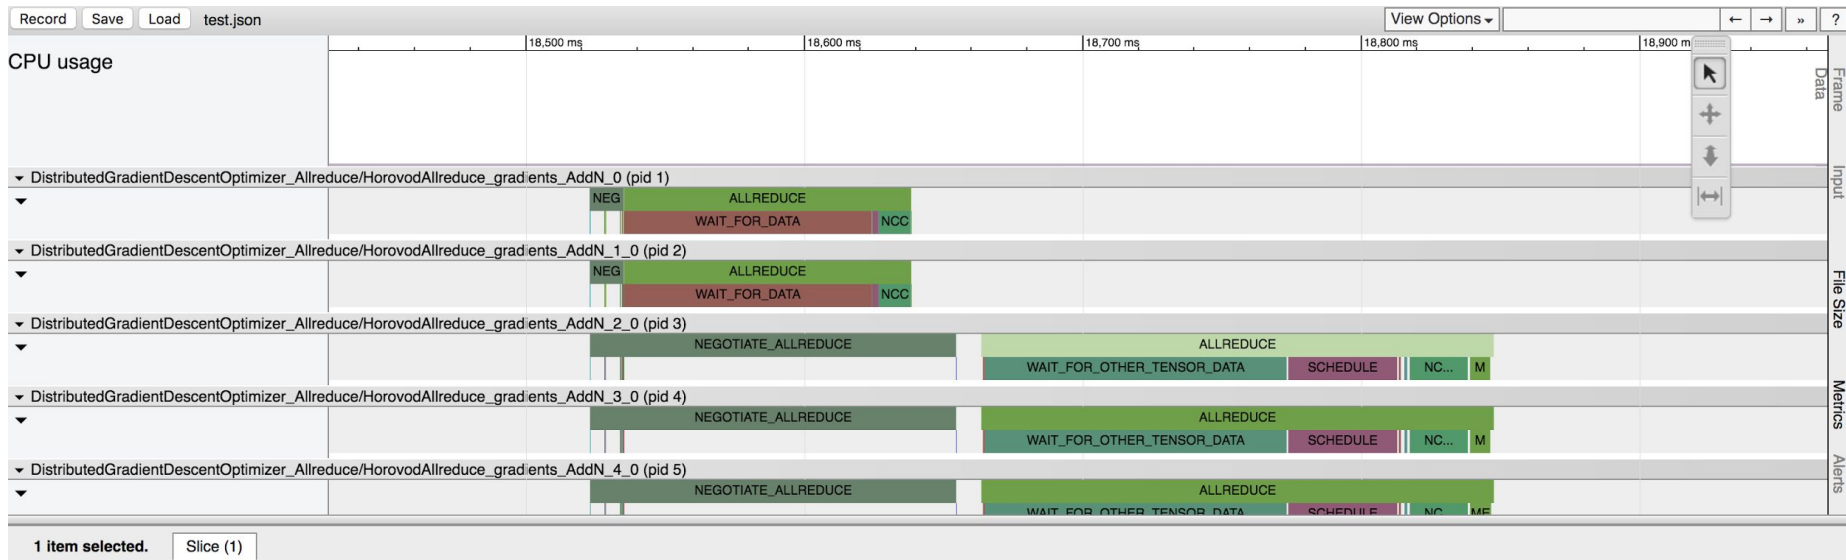
- ```
$ mpirun -np 4 \
 -H localhost:4 \
 -bind-to none -map-by slot \
 -x NCCL_DEBUG=INFO -x LD_LIBRARY_PATH \
 python train.py
```

- Run on 4 machines with 4 GPUs (Open MPI 3.0.0):

- ```
$ mpirun -np 16 \  
    -H server1:4,server2:4,server3:4,server4:4 \  
    -bind-to none -map-by slot \  
    -x NCCL_DEBUG=INFO -x LD_LIBRARY_PATH \  
    python train.py
```


- Boilerplate `mpirun` arguments are easily hidden in a convenience script

Debugging - Horovod Timeline



- Discovered that ResNet-152 has a lot of tiny tensors
- Added Tensor Fusion - smart batching causes large gains (bigger gain on less optimized networks)


Horovod Performance



UBER


With Horovod, same ResNet-101 can be trained for one epoch on ImageNet in 1.5 minutes. Scaling efficiency is improved to 88%, making it twice as efficient as standard distributed TF.

Horovod Performance Cont.



Practical Aspects - Initialization


- Use broadcast operation to make sure all workers start with the same weights
- Otherwise, averaged gradient will not point towards minimum (shown in red)



Practical Aspects - Data Partitioning


- Shuffle the dataset
- Partition records among workers
- Train by sequentially reading the partition
- After epoch is done, reshuffle and partition again

NOTE: make sure that all partitions contain the same number of batches, otherwise the training will reach deadlock



Practical Aspects - Random Sampling

- Shuffle the dataset
- Train by randomly reading data from whole dataset
- After epoch is done, reshuffle



Practical Aspects - Data


- Random sampling may cause some records to be read multiple times in a single epoch, while others not read at all
- In practice, both approaches typically yield same results
- **Conclusion:** use the most convenient option for your case
- **Remember:** validation can also be distributed, but need to make sure to average validation results from all the workers when using learning rate schedules that depend on validation
 - Horovod comes with `MetricAverageCallback` for Keras

Practical Aspects - Learning Rate Adjustment

- Facebook in paper “[Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour](https://arxiv.org/abs/1706.02677)” (arxiv.org/abs/1706.02677)

recommends linear scaling of learning rate:

- $LR_N = LR_1 * N$
 - Requires smooth warmup during first K epochs, as shown below
 - Works up to batch size 8192
- Horovod comes with `LearningRateWarmupCallback` for Keras



Practical Aspects - Learning Rate Adjustment Cont.

- Yang You, Igor Gitman, Boris Ginsburg in paper “Large Batch Training of Convolutional Networks” demonstrated scaling to batch of 32K examples (arxiv.org/abs/1708.03888)
 - Use per-layer adaptive learning rate scaling
- Google published a paper “Don't Decay the Learning Rate, Increase the Batch Size” (arxiv.org/abs/1711.00489) arguing that typical learning rate decay can be replaced with an increase of the batch size

Practical Aspects - Checkpointing & Logs

- Typically, a server would have multiple GPUs
- To avoid clashes, write checkpoints, TensorBoard logs and other artifacts on worker 0:
 - `if hvd.rank() == 0:`
 `# write checkpoint`

Practical Results at Uber

- Used Facebook's learning rate adjustment technique
- Trained convolutional networks and LSTMs in hours instead of days or weeks with the same final accuracy
- You can do that, too!

Giving Back

Horovod is available on GitHub:

<https://github.com/uber/horovod>



Thank you!

Horovod on our Eng Blog: <https://eng.uber.com/horovod>

Michelangelo on our Eng Blog: <https://eng.uber.com/michelangelo>

ML at Uber on YouTube: <http://t.uber.com/ml-meetup>



UBER

Proprietary and confidential © 2017 Uber Technologies, Inc. All rights reserved. No part of this document may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval systems, without permission in writing from Uber. This document is intended only for the use of the individual or entity to whom it is addressed and contains information that is privileged, confidential or otherwise exempt from disclosure under applicable law. All recipients of this document are notified that the information contained herein includes proprietary and confidential information of Uber, and recipient may not make use of, disseminate, or in any way disclose this document or any of the enclosed information to any person other than employees of addressee to the extent necessary for consultations with authorized personnel of Uber.